

Dynamic Lock Violation for Fault-tolerant Distributed Database System

Hua Guo

School of Information
Renmin University of China
Beijing, China
guohua2016@ruc.edu.cn

Xuan Zhou

School of Data Science And Engineering
East China Normal University
Shanghai, China
xzhou@dase.ecnu.edu.cn

Abstract—Many modern cloud distributed database management system (DBMS) scale horizontally by sharding its data on many nodes for scalability. Most of the databases in this category also build their transactional layer upon a replication layer for fault-tolerance. The replication layer uses a consensus protocol to reach consistency and implement automatic fault recovery. A transaction takes less time to enforce a serializable schedule than write its commit log to the replicated state machine(RSM). Thus, the replication layer and consensus protocol amplify transactions' lock duration. Exploit speculative techniques, such as controlled locked violation(CLV) and early lock release(ELR) can shorten lock duration, can optimize transaction performance, especially handle a high degree of contention. However, these techniques, which are mainly focused on single-site database and failed to scale achieve both performance and correctness on a distributed environment. In this paper, we introduce dynamic lock violation(DLV) which we designed for the distributed transaction, especially which is on a geo-replication layer for fault-tolerance. DLV can violate lock at a proper time to get the best performance and achieve both performance and correctness.

Index Terms—Database System, Distributed Transaction, Locking, High Availability

I. INTRODUCTION

Modern cloud distributed database scale-out by partitioning data into multiple nodes, so it can run transactions on different servers in parallel and increase throughput. However, when the database needs to access multiple partitions, it uses a coordinate protocol to ensure a transaction's atomicity. Distributed transactions usually lead to significant performance degradation, mainly due to the following reasons [1]:

1. Coordinating to commit needs a chatty protocol (i.e., two-phase commit) which causes more message overhead;
2. The message transmission overlaps with the critical path of transaction commit, which worsens the contention among transactions.

Furthermore, distributed DBMS in this category also use a replication layer below the transaction layer to guarantee fault tolerance. Transactional layer uses a specific concurrency control(CC) scheme to enforce a serializable schedule and a distributed commit protocol if transaction access multiple shards. The replicated layer often uses a Paxos-like consensus protocol to guarantee data replicas consistency. Typical implementation optimized replication performance by splitting data into very small chunks and build replicated state machines on them.

Although building multiple replicated state machine improve replication performance, it makes distributed transaction even more inevitable, as distributed transactions can more likely occur on different chunks of data.

Figure 2 presents the typical architecture. First, the database partitions its data with many shards to scale. Second, Each shard works on a replication layer which replicated data in several availability zones(AZ) [2] for high availability. Between different available zones, the replication layer uses a consensus protocol to shield consistency.

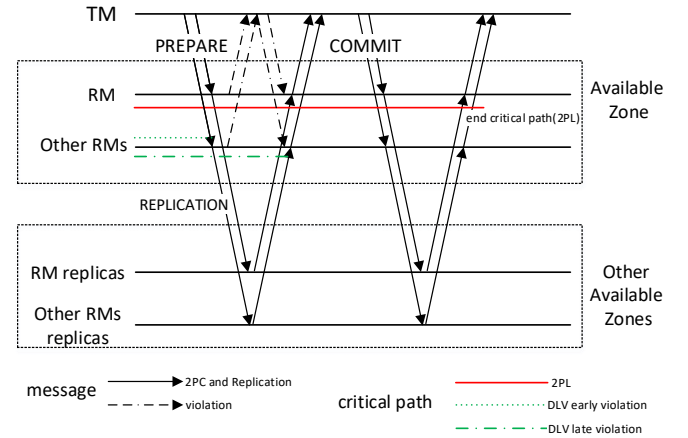


Fig. 1: The message flow and lock holding critical path of DB who uses 2PL(or with DLV) and 2PC works on replication layer. The dash arrow line message is introduced by DLV. The red lines show the critical path of S2PL and the green dash line shows the critical path of DLV.

This architecture uses a chatty message protocol fails to scale high contention workload, as much previous work has discussed [1] [3] [4]. But this architecture supports a wide range of transaction models and runs well on many workloads. Many industry distributed DBMS choose this two-layer architecture, such as Google Spanner [5] [6], NuoDB [7], CockroachDB [8], TiDB [9].

The distributed commit and replication coordination protocol enlarge the timespan of the critical path and amplified contention cost. We focus on distributed DBMS which

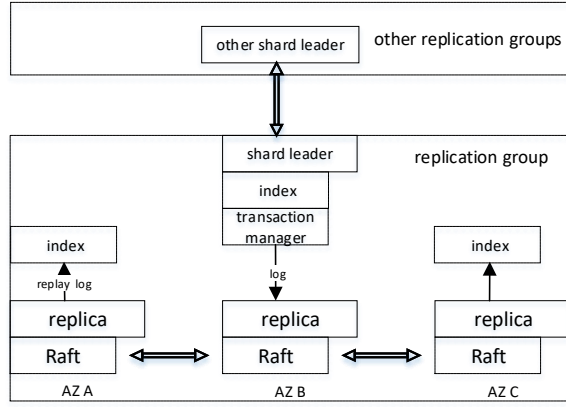


Fig. 2: Distributed and Replicated DBMS Architecture

uses locking scheme and coordinate protocol on a replicated layer, especially who running transaction processing on geo-replicated layer. Figure 1 shows the message flow of a distributed transaction using S2PL and 2PC works on a WAN replication layer. When a transaction requests its commit, *TM*(*transaction manager*) issues a prepare message to each *RM*(*resource manager*). *RM* replicates its result to other replicas to make it fault tolerant. *RM* responses its decision to *TM* after reach a consensus to guarantee fault tolerance. *TM* then collects all the *RM*s decisions. And then it issues either a commit or abort decision to *TM*'s replicas and broadcast the final result to all *RM*s. Once a *RM* receives the final result from *TM*, it can release the locks that it had retained since it first accesses the specified tuple. We depict the lock duration with the red line in Figure 1 when commit. The lock duration covers many message round trips including those over WAN. Locking for generating a serializable order of concurrency operations is can be overly lengthy to commit contented transactions. Such a commit and replication protocol will severely impair the concurrency when confronted with a high degree of contention.

Previous works used speculative techniques, such as early lock release(ELR) [10], controlled lock violation (CLV) [11] to optimize transaction processing using locking. These techniques can be extended to a distributed environment to improve concurrency. The two-layer architecture shares the same bottleneck with single node DBMS on forcing transaction log and faces even worse conditions. Distributed transactions work on geo-replicated DBMS need more time to write a log than non-distributed and non-replicated one. However, extended these techniques on the distributed environment is complex. There are some design considerations to choose from. To combine the two-phase commit protocol(2PL), violate(or release) lock at which phase transaction? As previous work addressed [11], violating lock at the first phase can exploit more concurrency but takes more dependency tracing cost and cascade abort cost. Violating lock at the second phase need to maintain less dependency and get less cascade abort rate. But it may not get

better concurrency. Transaction models, interactive or one-shot transactions, may have different message flow, how different transaction types could benefit from these techniques? Not all the transactions can benefit from CLV or ELR, little conflicts workloads are such cases. In this paper, we propose a dynamic lock violation(DLV) to boost the locking-based distributed DBMS, especially who are geo-replicated with a high commit latency. DLV decides lock violation time by runtime statistic. It maintains less commit dependencies and bears less cascade abort penalty compares previous implementation [11]. This section is the overall introduction of this paper. Section II is a review of related work. Section III presents a strict schedule is not necessary and hurt the performance of distributed and replicated DBMS. Section IV introduces DLV's implementation. Section V evaluates DLV and compare it with previous work. Section VI draws the conclusion of this paper.

II. RELATE WORK

This section introduces the related work of this paper.

A. Distributed Transaction on Replicated Layer

Recently, there are many scalable DBMS arisen in both academia and industry. Most of the systems in this category supports distributed query processing and replicate data across several data center geo-located in different areas for fault tolerance. A fault-tolerant database relies on state-machine replication(SMR) log to avoid single point failure. SMR needs to use a consensus protocol to enforce the same order of different replicas. Paxos [12] [13] is the most well-known consensus protocol. Paxos use two messages round trip to accept a value, one roundtrip for choosing a proposal and another to propose the value. Multidecree Paxos [14] elects a leader as the only proposer to eliminate Paxos first message roundtrip during normal processing. Raft [15] is a similar consensus protocol to Paxos, which is designed for understandability. Google spanner [5] [6] is a geo-replicated and shared-nothing DBMS that uses hardware clock for timestamp generation. VoltDB [16] is a main memory database who runs single threaded execution per partition. [1] use a deterministic transaction model, Calvin can commit distributed transaction without coordination protocol. VolteDB and Calvin, by using deterministic scheduling, they can use active replication to replicate transaction input rather than transaction effect. Consensus introduces significant overhead for its lots of message round trips and heavy network traffic. Tapir [3] use inconsistent replication layer and build consistent transaction on it to guarantee user level consistent. Janus [4] gets fewer wide-area round trips by consolidating the concurrency control and consensus and use deterministic serializable graph tracing to commit transactions under conflicts. Tapir and Janus, which benefit from their codesign of transaction and replication layer, can commit a distributed transaction in only one wide-area round trips.

B. Locking Concurrency Control

DBMS use concurrency control(CC) to calculate a serializable schedule for concurrent transactions. Two-phase lock-

ing(2PL) is the most widely used CC scheme. As a pessimistic method, 2PL assumes that it is likely that transactions will conflict. 2PL uses a lock to enforce the order of conflicting transactions. Strict 2PL(S2PL), in addition to 2PL, preserves its lock until a transaction's termination. S2PL guarantees transaction's recoverability but a 2PL schedule cannot. For enabling a simple recovery algorithm, most locking based databases choose S2PL. When extending S2PL to distributed databases, S2PL can take more time blocking on its commit critical path for additional message round trips.

2PL protocol implementation varies on how to process deadlock. In *no-wait* [17] policy, a transaction would immediately abort the transaction if it fails to lock record. Previous work has prove it is the most scalable technique to handle locking scheme, even in distributed environment [18] [17]. Another policy is *wait-die* [19] which is similar to *no-wait*. Transactions avoid false-positive aborting base on their start timestamps when database using 2PL *wait-die*. In *deadlock detection* [20], transactions can wait for each other without controlling. The transaction would abort only if there is a deadlock. *Deadlock detection* detect deadlock by explicitly tracing wait-for graph and testing circles. Many traditional single node database [21] [22] use *deadlock detection* technique because it has no false positive abort. Deadlock detection on distributed DBMS requires substantial network messages to identify circles and is costly.

C. Exploit Speculation and Lock Violation

Exploit speculation is not a new idea. Similar approaches have been introduced by many previous works. Early lock release (ELR) [23] [24] [25] [10] [26] shares the same idea with speculative approach. ELR can release transactions' lock without waiting for commit record flushed to disk. DeWitt et al. [23] firstly described ELR without implementation. Soisalon-Soininen et al. [24] proved that the correctness of ELR. Johnson et al. [25] and [10] evaluated the performance improvement made by ELR. Kimura et al. [10] [25] also address the weakness of previous ELR implementation [23] can produce wrong results for read-only transactions. Previous work exploits speculation mostly designed for single machine database system [24] [25] [10]. Jones et al. [27] use a restrict transaction model [28] implement speculation. Control lock violation (CLV) [11] achieve the same performance as ELR but with a simple and general implementation. CLV can apply to distributed databases and optimize both phases of two-phase commit. CLV can use a "register and report approach (RARA)" [29] to implement its dependency. RARA work well on a single-site database. When RARA is used to process a distributed transaction, the dependency tracing may be complex and costly. More cascade abort rates on distributed transaction also lead to more false-positive violations and carry a performance penalty.

III. BEYOND STRICT SCHEDULE AND LOCK VIOLATION

In the following subsections, we describe our preliminaries and assumptions, the basic rule to keep transaction correctness

when violating lock.

A. Preliminaries and Assumptions

The database shards its data by database primary keys. Each shard replicated its physiological logs across different AZs for fault-tolerance. The physiological logs record the row-level write operations of each transaction. The replicated layer use Raft protocol to sustain consensus of log order. Other replication protocols may also work. Only replica leader processes the transactional operations. Both one-shot and interactive transactions are supported.

Figure 3 shows the message flow of committing these two different type of transactions. As shown in these figures, The interactive transaction needs more message roundtrips compared to one-shot one. These two types of transaction model employ different speculative timing, which we will explain subsequently.

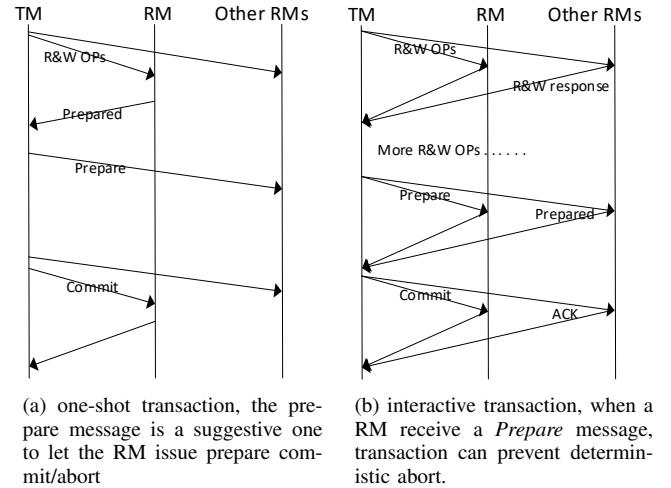


Fig. 3: Commit message flow of one-shot and interactive distributed transaction

B. Strict Scheduler is Too "Strict" for Correctness

Before we develop our method, we firstly review the formal definition of the transaction processing operation. Given a distributed transaction T_i , it runs on m sites $S = \{s_1, s_2, \dots, s_m\}$. The transaction history is a collection $H = \{h_1, h_2, \dots, h_m\}$, in which $h_u (1 \leq u \leq m) = \Pi_u(H)$ is the local history on site s_u . $\Pi_u(H)$ is H 's projection on site s_u . For any projected history $h_u (1 \leq u \leq m)$, h_u of transaction T_i includes a collection of operations o_i . o_i can be read or write operation, or command operations include prepare commit (abort), abort or commit. $r_i[x]$ donates transaction T_i reads record x , $w_i[x]$ donates transaction T_i writes record x . c_i , a_i , p_i^c , p_i^a mean transaction T_i commits, aborts, prepares commit or prepare abort respectively. Transaction abort due to many reasons, they can be: 1. User request abort, include access on some non-exists records; 2. Violation of serializability; 3. Database node crash for failure. We call the first two abort reasons *deterministic abort* and the last *non-deterministic abort*.

Transaction T_j has a *commit dependency* on transaction T_i , written as $T_j \rightarrow T_i$, if T_j can commit only if T_i commit. If T_j aborts, T_j need also abort. There are three kinds of dependencies, *wr-dependency*, *ww-dependency* and *rw-dependency*.

If transaction T_i and T_j have direct write-read conflict on record x , in any local history h , T_j read T_i 's write record, we call this dependency read-write(wr) dependency and denoted by $w_i[x] \rightarrow r_j[x]$.

Similarly, if T_i and T_j have direct write-write conflict on record x , T_j overwrite T_i 's write record, this is write-write(ww) dependency and written as $w_i[x] \rightarrow w_j[x]$.

And if T_i and T_j have direct read-write conflict on record x , T_j write x after T_i reads x , it is a read-write(rw) dependency and recorded as $r_i[x] \rightarrow w_j[x]$. A transaction T_j *speculative access* a record x , if there is another transaction T_i , T_j has a commit dependency on T_i and T_j access x before T_i has committed. We write this *danger dependency* as $w_j[x] \rightarrow_s r_i[x]$, $w_j[x] \rightarrow_s w_i[x]$, $r_j[x] \rightarrow_s w_i[x]$. We also write $T_j \rightarrow_s T_i$ to indicate transaction T_i has a commit *danger dependency* on T_j .

Traditional transaction schedulers choose strictness [30] to simplify implementation and avoid expensive transaction recovery cost. Strictness implies that a transaction cannot read or overwrite a previous write by another transaction which has not ended yet. For a locked base concurrency control scheme, the lock will hold until the transaction end, namely strict two-phase locking(S2PL). Strictness is not necessary to produce a correct schedule.

In schedule H_1 of Figure 4, there are 3 transactions working 3 shards, S_1 , S_2 , S_3 . There are dependencies, $r_1[x] \rightarrow w_3[x]$, $w_1[x] \rightarrow r_2[y]$, $r_2[y] \rightarrow w_3[x]$ and there is $T_1 \rightarrow T_2 \rightarrow T_3$. There is no circle in this dependency graph and the schedule is serializable and strict. Figure 5 shows an example of a non-strict but correct schedule H_2 . There is three records x , y , z , located at shard S_1 , S_2 , S_3 . Transaction T_1 execute write y , write x . Transaction T_2 read T_1 's write on x before T_1 commits. Transaction T_3 overwrite T_2 's write ahead T_2 's commit. The history

$$H = \{h_1, h_2, h_3\},$$

in which,

$$\begin{aligned} h_1 &= w_1[x]w_3[x]p_1^c c_1 p_3^c c_3 \\ h_2 &= w_1[y]r_2[y]p_1^c c_1 p_2^c c_2 \\ h_3 &= r_2[z]w_3[z]p_2^c c_2 p_3^c c_3 \end{aligned}$$

is not a strict history, but it's a serializable history. Both of these two schedules are serializable equate with the serial schedule, T_1 , T_2 and T_3 . Both schedule H_1 and H_2 are correct.

To get better concurrency, the scheduler can produce schedule like H_2 in Figure 5. Suppose locking schedule H_2 , then transaction T_1 must release its locks or make its locks can be violative before it knows its commit decision. A transaction commits by no means an operation in a flash but progress that needs take lots of time, especially when it is a distributed commit on an RSM. Strictness scheduler on a distributed and replicated DBMS has a long critical path. Our basic idea

S_1	$w_1[x]$	pc_1	c_1	$w_3[x]$	pc_3	c_3
S_2	$w_1[y]$	pc_1	c_1	$r_2[y]$	pc_2	c_2
S_3				$r_2[z]$	pc_2	c_2
				$w_3[z]$	pc_3	c_3

Fig. 4: A strict and serializable schedule H_1

S_1	$w_1[x]$	$w_3[x]$	pc_1	c_1	pc_3	c_3
S_2	$w_1[y]$	$r_2[y]$	pc_1	c_1	pc_2	c_2
S_3	$r_2[z]$	$w_3[z]$	pc_2	c_2	pc_3	c_3

Fig. 5: A non-strict but serializable schedule H_2

is to develop a serializable but non-strict correct scheduler for distributed transactions and shorten the critical path when commit. A single node transaction can exploit log order to maintain dependencies because dependent transactions write their logs orderly [23] [10]. When we extended the non-strict locking scheme to a distributed transaction, transaction dependency maintaining is more complex. The serializable scheduler is a correct one if it can prevent all aborted transaction. But when there are aborted transactions, it does not. To produce the correct log, the schedule must be both commit serializable and recoverable [31]. Transactions need to maintain commit dependencies to guarantee serializability and recoverability when exploiting non-strictness.

C. Lock Violating Rules and Dependency Tracing

Lock violation scheduler must follow some rules to preserve correctness. First, a lock violation scheduler should create serializable schedules. Consider a schedule H_3 in Figure 6 as an example. There are danger dependencies in H_3 :

$$w_3[x] \rightarrow_s w_1[x], w_1[y] \rightarrow_s r_2[y], r_2[z] \rightarrow_s w_3[z]$$

S_1	$w_3[x]$	$w_1[x]$	pa_1	a_1
S_2	$w_1[y]$	$r_2[y]$	pc_1	a_1
S_3	$r_2[z]$	$w_3[z]$	pc_2	

Fig. 6: Schedule H_3 , T_1 abort due to non-serializable

S_1	$w_1[x]$	$w_3[x]$	pa_1	a_1
S_2	$w_1[y]$	$r_2[y]$	pc_2	c_2
S_3	$r_2[z]$	$w_3[z]$	pc_2	c_2

Fig. 7: Schedule H_4 , T_2 commit ahead T_1 , non-recoverable anomaly

There is a circle $T_1 \rightarrow T_2 \rightarrow T_3 \rightarrow T_1$. This schedule is not serializable. Transaction T_2 read from uncommitted transaction T_3 . When T_1 abort for non-serializable, T_2 must cascade abort to avoid anomaly. Assume a lock violation scheduler creates schedule H_3 . Not formally discussing, a lock violation operation is just as a transaction releases its lock and another later transaction acquire locks. Then the schedule H_3 failed to comply with two phases principle.

For a traditional locking CC scheme, transaction T_2 needs to wait for T_1 's release its lock until T_1 commit. Lock violation scheduler guarantees the dependency graph has no circles by tracing commit dependencies if a transaction violates locks of another conflict transaction operation.

Secondly, a schedule generated by lock violation must also be recoverable. In schedule H_4 of Figure 7, T_2 read from uncommitted transaction T_1 's write. There is a danger dependency, $w_1[y] \rightarrow_s r_2[y]$ and T_2 's commit is ahead T_1 's commit. Schedule H_4 is non-recoverable. If T_1 abort, then T_2 will return an error y value. The scheduler requires to maintain dependencies to preserve the recoverability of the schedule.

Additionally, there are several design considerations and choices arisen except the correctness. Can lock violation easily adapt to different transaction models and 2PC? Violating lock at the first phase of 2PC may be superior to violating lock at the second phase because it can shorten the more critical path length. But the second phase violation may bear more cascade aborts which are useless work. DLV permits lock violation at two points in the timeline of running transactions. We call these *early violation* and *late violation*.

Suppose H is a schedule which is created by lock violation, then H is extended by adding the lock, unlock and lock violation operations. We write H by:

$$H = l_i[x]o_i[x]...vl_j[x]o_j[x]...l_i[y]o_i[y]...ul_i[x]...$$

In H , x , y are the same records. There are following lock/unlock/violation operations, transaction T_i locks record x ; T_j violating T_i 's lock on x ; T_j locks tuple y ; T_i release locks on x . If there is such $vl_j[x]$ and $l_i[y]$ operations in schedule H , then transaction T_j is *early lock violate* T_i 's lock on x ; otherwise, this is *late lock violation*. A scheduler uses *early lock violation* may lead to non-serializable schedule. Then DLV needs maintains all wr , ww and rw dependencies after violating locks and guarantee the dependency graph of the schedule is acyclic if using *early lock violation*. On the contrary, *late lock violation* cannot make an acyclic dependency graph to become a cyclic one by adding any dependency edges. This can be proved by formulating *late lock violation* as 2PL proving.

Assume that there is a wr -dependency from T_i to T_j . T_j , which can be written as $w_i[x] \rightarrow r_j[x]$. T_j cannot commit if T_i has not committed. Traditional S2PL schedule can guarantee this by release locking when T_i commit. Lock violating violates locking rule and T_j can read T_i 's write on x before T_i commits. In a lock violation schedule case, transactions must trace dependencies and commit as dependency orders. Composite with 2PC protocol, we have the following rules:

- 1) T_i prepares only if T_j commit;
- 2) T_i commits only if T_j commit;
- 3) If T_j aborts, T_i must also abort

By tracing dependencies after violating a lock, DLV schedule achieves both serializability and recoverability.

IV. DLV IMPLEMENTATION

In this section, we introduce DLV implementation. The following contents would include: How DLV can avoid complex recovery algorithm and maintain the most limited amount of dependencies; How DLV choose the proper time of enabling violation; The wait-die policy of DLV use; The pseudocode code description finally.

A. In Memory Speculative Versions

The non-strict scheduler needs more complex recovery algorithm to keep the correctness. Take a schedule H_5 as an example,

$$H_5 = w_1[x]w_1[y]r_2[x]w_2[y]a_1a_2$$

If transaction T_1 abort, this cause cascade abort. Traditional database use undo log to process recovery transaction write operations. Implementation undo log maybe a little bit tricky when exploiting non-strict. A wrong recovery expand schedule of T_1 may be like $exp(H_5)$, in which $w_i^-[x]$ means transaction T_i undo its write on x .

$$exp(H_5) = w_1[x]w_1[y]r_2[x]w_2[y] w_1^-[y]w_1^-[x]c_1w_2^-[y]c_2$$

Suppose the initial value of records x and y of are both 0. The value of records x , y and the undo log formatted after executing every operations in $exp(H)$ is shown in Table I. Finally, after the execution of this schedule, both transaction T_1 and T_2 aborts. The value of y is 1, which the correct result should be the initial value 0.

operations	x	y	undo
$w_1[x = 1]$	1	0	x=0
$w_1[y = 1]$	1	1	y=0
$r_2[x]$	1	1	
$w_2[y = 2]$	1	2	y=1
$w_1^-[y = 0]$	1	0	
$w_1^-[x = 0]$	0	0	
c_1	0	0	
$w_2^-[y = 1]$	0	1	
c_2	0	1	

TABLE I: x, y values, undo log after the execution of $exp(H_5)$

To tackle this anomaly, recovery must use a more complex algorithm such as SOT [31]. For schedule H_5 , a correct recovery expansion may be:

$$exp^*(H_5) = w_1[x]w_1[y]r_2[x]w_2[y]w_2^-[y]c_2w_1^-[y]w_1^-[x]c_1$$

The scheduler must recovery transaction by the reserve order of write operation. If x and y is on the same database node and use *late lock violation*, the recovery of a transaction is simple. Because there is no partial failure, the transaction would commit in log order. No additional work is needed when system recovery using traditional Aries algorithm [32].

However, if x and y are not located at the same node, this undo operation order is hard to accomplish because of partial failure. When using *early lock violation*, there are similar problems since transaction recovery must also undo transactional operations by reserve order. To avoid this complexity, DLV maintains uncommitted speculative versions in memory and accepts no-steal policy when writing data. No-steal policy need storage cannot write uncommitted data to permanent storages. For most transactions would write a little data except the bulk loading ones and the modern database runs on a machine with large RAM, using no-steal policy to save memory is not necessary. By no-steal and speculative versions, the database needs no undo log, transaction rollback and failure recovery would be more simple and efficient. DLV's speculative version implementation is a little similar with many multi-version concurrency control scheme. The list is structured from the newest version to the oldest version and the last version of this list is the committed version. Speculation versions are always stored in main memory and needs no persistence. If a transaction would abort, it only needs to remove its write versions from speculative version list.

Previously, we have discussed that a ww dependency does not affect recoverability. *Late lock violation*, since it has promised serializability, so it can ignore ww and rw dependencies and only trace wr dependencies for recoverability. Figure 8 show a series of schedule access on two contention rows, x , y . The green rectangles are speculative versions and the red ones are committed versions. Although there is ww dependency $w_6[x] \rightarrow w_4[x]$. The abort of T_4 does not cause T_6 cascade abort.




Fig. 8: speculative(green) and committed(red) versions, x_i , y_i express this version is from transaction T_i 's write

B. Dynamic Decide Early or Late Violation

Early-violation can be more appropriate than *late lock violation* when there are less cascade abort caused by non-

deterministic abort. *Early-violation* also needs more dependency tracing costs. Too many cascade abort can lead to a lot of useless work.

We implement *late lock violation* by adding a message round trip to prevent deterministic abort. This additional message flow shows in Algorithm 4. In Figure 1, the message is shown as dotted arrow lines. Before an RM decides to replicate its prepare log, it also sends a *Ready* message to TM and tells TM it will prepare this transaction. *Ready* message shows that the RM will prepare commit or prepare abort. When the TM collects all RM 's *Ready* message, it sends *Violate* messages to tell every RM s make their locks violatable. An interactive-transaction can combine these messages with the last operations and prepare requests in passing, as Figure 3 shows. For a one-shot transaction, this message flow also takes less time than the overall message flow of 2PC because of no log replication time delay. Especially when all the RM s and the TM are located in LAN, there is no WAN message needed.

DLV records transaction statistic information to decide use *early lock violation* or *late lock violation*. We define a distributed transaction working on one RM as partial transaction and a partial transaction enters prepared commit phase but failed to commit finally as partial prepare. DLV calculates partial prepare rate at a period to decide which violation strategy to choose. We suppose that a transaction T running at time τ would access a collection of shards S_1, S_2, \dots, S_n . The message-round trip time from T 's TM to RM on shard S_i is RTT_i . In a window period time from $\tau - \delta$ to τ , there is N_p partial prepares of total N_t partial transactions. DLV would test the following conditions where Θ is a constant coefficient.

$$N_p/N_t < \Theta * \max(RTT_i)$$

DLV would choose *early lock violation* if this condition is satisfied. Otherwise, it would use *late lock violation*.

C. Locking Violation and Maintain Commit Dependency

DLV uses *wait-die* protocol to avoid deadlock. At the beginning of a transaction, the transaction uses the current timestamp to generate a transaction id. The conflict transaction operations are queued based on their transaction id's order.

DLV uses register and report [29] to maintain dependencies. Every transaction context stores an in-dependency transaction(in_dn) number count to record how many transactions it is dependent on. The transaction also keeps an out-transaction set(out_set) attribute to record the transactions it depends on. When a transaction T speculatively reads from S , T registers dependency from S by adding T to out_set of S and increasing in_dn of T by one. These steps are described by lines 4 - 7 in function READ in Algorithm 1. A transaction cannot prepare if its in_dn value is greater than 0, which means some in-dependency transaction does not commit yet. If a transaction's in_dn value is less than 0, the transaction must abort because there is some in-dependency abort causing a cascade abort.

Algorithm 2 shows how to prepare a transaction. When a transaction commits, this transaction would traversal its *out_set* and decrease every transaction's *in_dn* by one, this is shown in line 6 - 11 of Algorithm 3 function. If a transaction aborts, it may cause a cascade abort. Line 2 of function CASCADE 16 shows the assign *in_dn* by a negative value when cascade abort.

D. Pseudocode Description

Algorithm 1 shows the execution phase of a transaction. Algorithm 2 shows the prepare phase of a transaction. Algorithm 3 shows the commit phase of a transaction. Algorithm 4 shows the speculation phase of a transaction.

Algorithm 1 Execution phase of transaction T . Read and write a key

```

1: function READ( $T, key$ )
2:    $newest\_version \leftarrow Head(Tuple(key).version\_list)$ 
3:   if  $newest\_version$  is created by transaction  $S$ 
   and  $key$  is ICommit locked by  $S$  then
4:     if  $T \notin S.out\_set$  then
5:        $S.out\_set \leftarrow S.out\_set \cup T$ 
6:        $T.in\_dn \leftarrow T.in\_dn + 1$ 
7:     end if
8:   end if
9:   if  $key$  is write locked by transaction  $S$  then
10:     $S.wait \leftarrow S.wait + 1$ 
11:    wait lock till die
12:    if die then
13:       $T.no\_da \leftarrow \text{False}$ 
14:      return die error.
15:    end if
16:   end if
17:   if  $key$  is IAbort locked by transaction  $S$  then
18:    wait lock this lock released
19:   end if
20:   Lock( $T, key, Read$ )
21:   return  $key$ 's value.
22: end function

1: function WRITE( $T, key, value$ )
2:   if  $key$  is read or write locked then
3:     if  $key$  is write locked by transaction  $S$  then
4:        $S.wait \leftarrow S.wait + 1$ 
5:     end if
6:     wait lock till die
7:     if die then
8:        $T.no\_da \leftarrow \text{False}$ 
9:       return die error.
10:    end if
11:   end if
12:   Lock( $T, key, Write$ )
13:   add a new version of  $key$ 's tuple, assign  $value$ 
14: end function

```

Algorithm 2 Prepare phase of transaction T

```

1: function PREPARE( $T$ )
2:   wait if  $T.in\_dn > 0$ 
3:   if  $T.in\_dn < 0$  then
4:     response  $TM$  message {Prepare Abort}
5:   else if  $T.in\_dn = 0$  then
6:     response  $TM$  message {Prepare Commit}
   or {Prepare Abort}
7:   end if
8: end function

```

Algorithm 3 Commit phase of transaction T , commit and (cascade)abort function

```

1: function COMMIT( $T$ )
2:   garbage collect old version in
    $Tuple(key).version\_list$ 
3:   for  $key \in T.write\_set \cup T.read\_set$  do
4:     Unlock( $T, key, Read/Write$ )
5:   end for
6:   for  $T_{out} \in T.out\_set$  do
7:      $T_{out}.in\_dn \leftarrow T_{out}.in\_dn - 1$     ▷ keep exactly
   once
8:     if  $T_{out}.in\_dn = 0$  then
9:       report  $T_{out}.in\_dn = 0$     ▷ stop waiting on
   function PREPARE line 2
10:    end if
11:   end for
12:   response  $TM$  message {Commit ACK}
13: end function

1: function ABORT( $T$ )
2:   call CASCADE( $T$ )
3:   for  $key \in T.write\_set \cup T.read\_set$  do
4:     Unlock( $T, key, Read/Write$ )
5:   end for
6:   response  $TM$  message {Abort ACK}
7: end function

1: function CASCADE( $T$ )
2:    $T.in\_dn \leftarrow -\infty$ 
3:   for  $key \in T.write\_set$  do
4:     if  $key$  is ICommit locked by  $T$  then
5:       ModifyLock( $T, key, IAbort$ )
6:     end if
7:     for  $version \in Tuple(key).version\_list$  do
8:       if  $version$  is created by  $T$  then
9:         remove  $version$  from list
10:        break
11:      end if
12:    end for
13:   end for
14:   for  $T_{out} \in T.out\_set$  do
15:     call CASCADE( $T_{out}$ )
16:   end for
17: end function

```

Algorithm 4 Speculate phase.

Ready and *Speculate* works on *RM*.

TM call *Decide* send when *TM* collects all *RM*'s *Ready* message.

msgs is a collection of *Ready* message which *TM* receives from all the *RMs*.

Θ is a threshold value to enable speculation.

```

1: function READY(T)
2:   response TM message
   {Ready, wait  $\leftarrow$  T.wait, non_da  $\leftarrow$  T.non_da}
3: end function
1: function DECIDE(T, msgs)
2:   if  $\forall m \in msgs, m.non\_da$  is True and
    $\exists m \in msgs, m.wait > \Theta$  then
3:     send all RMs message {Speculate}
4:   end if
5: end function
1: function SPECULATE(T)
2:   for key  $\in$  T.read_set do
3:     Unlock(T, key, Read)
4:   end for
5:   for key  $\in$  T.write_set do
6:     if key is Write locked by T then
7:       ModifyLock(T, key, ICommit)
8:     end if
9:   end for
10: end function

```

V. EXPERIMENTS AND EVALUATIONS

We develop a replicated distributed DBMS demo and evaluate the performance of *DLV*. As a comparison with *DLV*, we also implement S2PL wait die(S2PL) scheme, CLV optimize both violate at the 1st phase(CLV1P) and 2nd phase(CLV2P).

A. Experiments Setting

Our experiments performed on a cluster of 12 Aliyun ecs.g6.large server. Each server has 2 virtual CPU with 2.5GHz clock speed, 8GB RAM, runs Ubuntu 18.04. The data is partitioned by 4 shards, each shard has 3 replicas which is replicated across 3 AZs, which is located at Heyuan, Hangzhou and Huhehot. Every AZ has a full data copy of each shard. The internal network bandwidth of each AZ is 1Gbps. We choose a modifies version TPCC and YCSB workload. All the transactions are distributed transactions. The TPCC data is sharded by the warehouse id. The Item table is replicated to all shards. Each transaction will retry after 3 seconds if it aborts for violation serializability. The evaluation both tested on both scattered (leader) mode and gathered (leader) mode. In gathered mode, all of the replica leaders are located in the same AZs. In scattered mode, the replica leaders are not located in the same AZs.

B. TPCC Performance Evaluation

Figure 9 shows the NewOrder performance of when adding terminal numbers of each node in the gathered mode (all the

TM and RMs are located in the same AZ) and scattered mode. (TMs and RMs are located in the different AZs).




Fig. 9: throughput and abort rate of different terminal number of each node when running TPCC NewOrder transactions in gathered mode




Fig. 10: throughput and abort rate of different warehouse number when running TPCC NewOrder transactions in gathered mode




Fig. 11: throughput and abort rate of different warehouse number when running TPCC Payment transactions in gathered mode

C. YCSB Performance Evaluation

VI. CONCLUSION

We extend *CLV* to distributed transaction and evaluate its performance on a replicated layer. Our distributed version *CLV* is *DLV*. It dynamically decides to violate lock at the most suitable time. According to our evaluation, *DLV* can improve performance of contention workload for shortening critical path. It avoids many unnecessary dependency tracing cost and cascade abort penalty against previous work.




Fig. 12: throughput and abort rate of different possible cascade abort rate, when running TPCC NewOrder transactions in gathered mode




Fig. 13: throughput and abort rate of different possible cascade abort rate, when running TPCC NewOrder transactions in scattered mode




Fig. 14: throughput and abort rate of different warehouse number when running YCSB transactions in gathered mode




Fig. 15: throughput and abort rate of different warehouse number when running YCSB transactions in scattered mode

- [1] A. Thomson, T. Diamond, S. Weng, K. Ren, P. Shao, and D. J. Abadi, "Calvin: fast distributed transactions for partitioned database systems," in *Proceedings of the ACM SIGMOD International Conference on Management of Data, SIGMOD 2012, Scottsdale, AZ, USA, May 20-24, 2012*, K. S. Candan, Y. Chen, R. T. Snodgrass, L. Gravano, and A. Fuxman, Eds. ACM, 2012, pp. 1–12. [Online]. Available: <https://doi.org/10.1145/2213836.2213838>
- [2] A. Verbitski, A. Gupta, D. Saha, J. Corey, K. Gupta, M. Brahmadesam, R. Mittal, S. Krishnamurthy, S. Maurice, T. Kharatishvili, and X. Bao, "Amazon aurora: On avoiding distributed consensus for i/os, commits, and membership changes," in *Proceedings of the 2018 International Conference on Management of Data, SIGMOD Conference 2018, Houston, TX, USA, June 10-15, 2018*, G. Das, C. M. Jermaine, and P. A. Bernstein, Eds. ACM, 2018, pp. 789–796. [Online]. Available: <https://doi.org/10.1145/3183713.3196937>
- [3] I. Zhang, N. K. Sharma, A. Szekeres, A. Krishnamurthy, and D. R. K. Ports, "Building consistent transactions with inconsistent replication," in *Proceedings of the 25th Symposium on Operating Systems Principles, SOSP 2015, Monterey, CA, USA, October 4-7, 2015*, E. L. Miller and S. Hand, Eds. ACM, 2015, pp. 263–278. [Online]. Available: <https://doi.org/10.1145/2815400.2815404>
- [4] S. Mu, L. Nelson, W. Lloyd, and J. Li, "Consolidating concurrency control and consensus for commits under conflicts," in *12th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2016, Savannah, GA, USA, November 2-4, 2016*, K. Keeton and T. Roscoe, Eds. USENIX Association, 2016, pp. 517–532. [Online]. Available: <https://www.usenix.org/conference/osdi16/technical-sessions/presentation/mu>
- [5] J. C. Corbett, J. Dean, M. Epstein, A. Fikes, C. Frost, J. J. Furman, S. Ghemawat, A. Gubarev, C. Heiser, P. Hochschild, W. C. Hsieh, S. Kanthak, E. Kogan, H. Li, A. Lloyd, S. Melnik, D. Mwaura, D. Nagle, S. Quinlan, R. Rao, L. Rolig, Y. Saito, M. Szymaniak, C. Taylor, R. Wang, and D. Woodford, "Spanner: Google's globally-distributed database," in *10th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2012, Hollywood, CA, USA, October 8-10, 2012*, C. Thekkath and A. Vahdat, Eds. USENIX Association, 2012, pp. 251–264. [Online]. Available: <https://www.usenix.org/conference/osdi12/technical-sessions/presentation/corbett>
- [6] D. F. Bacon, N. Bales, N. Bruno, B. F. Cooper, A. Dickinson, A. Fikes, C. Fraser, A. Gubarev, M. Joshi, E. Kogan, A. Lloyd, S. Melnik, R. Rao, D. Shue, C. Taylor, M. van der Holst, and D. Woodford, "Spanner: Becoming a SQL system," in *Proceedings of the 2017 ACM International Conference on Management of Data, SIGMOD Conference 2017, Chicago, IL, USA, May 14-19, 2017*, S. Salihoglu, W. Zhou, R. Chirkova, J. Yang, and D. Suciu, Eds. ACM, 2017, pp. 331–343. [Online]. Available: <https://doi.org/10.1145/3035918.3056103>
- [7] "Nuodb," <https://www.nuodb.com/>.
- [8] "Cockroachdb," [3] <https://www.cockroachlabs.com/>.
- [9] "Tidb," <https://pingcap.com/en/>.
- [10] H. Kimura, G. Graefe, and H. A. Kuno, "Efficient locking techniques for databases on modern hardware," in *International Workshop on Accelerating Data Management Systems Using Modern Processor and Storage Architectures - ADMS 2012, Istanbul, Turkey, August 27, 2012*, R. Bordawekar and C. A. Lang, Eds., 2012, pp. 1–12. [Online]. Available: http://www.adms-conf.org/kimura_adms12.pdf
- [11] G. Graefe, M. Lillibridge, H. A. Kuno, J. Tucek, and A. C. Veitch, "Controlled lock violation," in *Proceedings of the ACM SIGMOD International Conference on Management of Data, SIGMOD 2013, New York, NY, USA, June 22-27, 2013*, K. A. Ross, D. Srivastava, and D. Papadias, Eds. ACM, 2013, pp. 85–96. [Online]. Available: <https://doi.org/10.1145/2463676.2465325>
- [12] L. Lamport, "The part-time parliament," *ACM Trans. Comput. Syst.*, vol. 16, no. 2, pp. 133–169, 1998. [Online]. Available: <https://doi.org/10.1145/279227.279229>
- [13] —, "Paxos made simple, fast, and byzantine," in *Proceedings of the 6th International Conference on Principles of Distributed Systems. OPODIS 2002, Reims, France, December 11-13, 2002*, ser. Studia Informatica Universalis, A. Bui and H. Fouchal, Eds., vol. 3. Suger, Saint-Denis, rue Catulienne, France, 2002, pp. 7–9.

- [14] R. van Renesse and D. Altinbukan, "Paxos made moderately complex," *ACM Comput. Surv.*, vol. 47, no. 3, pp. 42:1–42:36, 2015. [Online]. Available: <https://doi.org/10.1145/2673577>
- [15] D. Ongaro and J. K. Ousterhout, "In search of an understandable consensus algorithm," in *2014 USENIX Annual Technical Conference, USENIX ATC '14, Philadelphia, PA, USA, June 19-20, 2014.*, G. Gibson and N. Zeldovich, Eds. USENIX Association, 2014, pp. 305–319. [Online]. Available: <https://www.usenix.org/conference/atc14/technical-sessions/presentation/ongaro>
- [16] "Voltdb," <http://www.voltdb.com/>.
- [17] R. Harding, D. V. Aken, A. Pavlo, and M. Stonebraker, "An evaluation of distributed concurrency control," *Proc. VLDB Endow.*, vol. 10, no. 5, pp. 553–564, 2017. [Online]. Available: <http://www.vldb.org/pvldb/vol10/p553-harding.pdf>
- [18] X. Yu, G. Bezerra, A. Pavlo, S. Devadas, and M. Stonebraker, "Staring into the abyss: An evaluation of concurrency control with one thousand cores," *Proc. VLDB Endow.*, vol. 8, no. 3, pp. 209–220, 2014. [Online]. Available: <http://www.vldb.org/pvldb/vol8/p209-yu.pdf>
- [19] P. A. Bernstein and N. Goodman, "Concurrency control in distributed database systems," *ACM Comput. Surv.*, vol. 13, no. 2, pp. 185–221, 1981. [Online]. Available: <https://doi.org/10.1145/356842.356846>
- [20] J. Gray, R. A. Lorie, G. R. Putzolu, and I. L. Traiger, "Granularity of locks and degrees of consistency in a shared data base," in *Modelling in Data Base Management Systems, Proceeding of the IFIP Working Conference on Modelling in Data Base Management Systems, Freudenstadt, Germany, January 5-8, 1976*, G. M. Nijssen, Ed. North-Holland, 1976, pp. 365–394.
- [21] "Mysql," <https://www.mysql.com/>.
- [22] "Postgresql," <https://www.postgresql.org/>.
- [23] D. J. DeWitt, R. H. Katz, F. Olken, L. D. Shapiro, M. R. Stonebraker, and D. Wood, "Implementation techniques for main memory database systems," in *Proceedings of the 1984 ACM SIGMOD international conference on Management of data - SIGMOD '84*. Boston, Massachusetts: ACM Press, 1984, p. 1. [Online]. Available: <http://portal.acm.org/citation.cfm?doid=602259.602261>
- [24] E. Soisalon-Soininen and T. Ylönen, "Partial strictness in two-phase locking," in *Database Theory - ICDT'95, 5th International Conference, Prague, Czech Republic, January 11-13, 1995, Proceedings*, ser. Lecture Notes in Computer Science, G. Gottlob and M. Y. Vardi, Eds., vol. 893. Springer, 1995, pp. 139–147. [Online]. Available: https://doi.org/10.1007/3-540-58907-4_12
- [25] R. Johnson, I. Pandis, R. Stoica, M. Athanassoulis, and A. Ailamaki, "Aether: A scalable approach to logging," *PVLDB*, vol. 3, no. 1, pp. 681–692, 2010. [Online]. Available: http://www.vldb.org/pvldb/vldb2010/pvldb_vol3/R61.pdf
- [26] P. A. Bernstein, "Actor-oriented database systems," in *34th IEEE International Conference on Data Engineering, ICDE 2018, Paris, France, April 16-19, 2018*. IEEE Computer Society, 2018, pp. 13–14. [Online]. Available: <https://doi.org/10.1109/ICDE.2018.00010>
- [27] E. P. C. Jones, D. J. Abadi, and S. Madden, "Low overhead concurrency control for partitioned main memory databases," in *Proceedings of the ACM SIGMOD International Conference on Management of Data, SIGMOD 2010, Indianapolis, Indiana, USA, June 6-10, 2010*, A. K. Elmagarmid and D. Agrawal, Eds. ACM, 2010, pp. 603–614. [Online]. Available: <https://doi.org/10.1145/1807167.1807233>
- [28] R. Kallman, H. Kimura, J. Natkins, A. Pavlo, A. Rasin, S. B. Zdonik, E. P. C. Jones, S. Madden, M. Stonebraker, Y. Zhang, J. Hugg, and D. J. Abadi, "H-store: a high-performance, distributed main memory transaction processing system," *PVLDB*, vol. 1, no. 2, pp. 1496–1499, 2008. [Online]. Available: <http://www.vldb.org/pvldb/1/1454211.pdf>
- [29] P. Larson, S. Blanas, C. Diaconu, C. Freedman, J. M. Patel, and M. Zwilling, "High-performance concurrency control mechanisms for main-memory databases," *PVLDB*, vol. 5, no. 4, pp. 298–309, 2011. [Online]. Available: http://vldb.org/pvldb/vol5/p298_perakelarsen_vldb2012.pdf
- [30] Y. Raz, "The principle of commitment ordering, or guaranteeing serializability in a heterogeneous environment of multiple autonomous resource managers using atomic commitment," in *18th International Conference on Very Large Data Bases, August 23-27, 1992, Vancouver, Canada, Proceedings.*, L. Yuan, Ed. Morgan Kaufmann, 1992, pp. 292–312. [Online]. Available: <http://www.vldb.org/conf/1992/P292.PDF>
- [31] G. Alonso, R. Vingralek, D. Agrawal, Y. Breitbart, A. El Abbadi, H. Schek, and G. Weikum, "Unifying concurrency control and recovery of transactions," *Inf. Syst.*, vol. 19, no. 1, pp. 101–115, 1994. [Online]. Available: [https://doi.org/10.1016/0306-4379\(94\)90029-9](https://doi.org/10.1016/0306-4379(94)90029-9)
- [32] C. Mohan, D. Haderle, B. G. Lindsay, H. Pirahesh, and P. M. Schwarz, "ARIES: A transaction recovery method supporting fine-granularity locking and partial rollbacks using write-ahead logging," *ACM Trans. Database Syst.*, vol. 17, no. 1, pp. 94–162, 1992. [Online]. Available: <https://doi.org/10.1145/128765.128770>